

# **CLOUD TAP: Enterprise Network Simulation System**

**Data Structures Course Project**

---

**Team Members:**

- Abdul Wahid (243699)
- Uneeb Ali (243622)

**Instructor:** Syd-Atta-Ur-Rehman

**Course:** Data Structures

**Submission Date:** December 29, 2025

# **Executive Summary**

**Problem:** Modern enterprise networks require efficient management of 48+ interconnected devices with dynamic resource allocation, path optimization, and real-time failure detection.

**Solution:** Cloud TAP (Tier Access Platform) - A C++ network simulation system leveraging advanced data structures to model realistic corporate infrastructure.

## **Key Achievements:**

- 48 active network devices across 5 departments
- 75 bidirectional physical connections
- 7 DHCP pools with automatic IP management
- $O(\log n)$  device lookup,  $O(V+E)$  graph traversal
- Real-time event tracking with 1000-entry circular buffer

**Technologies Implemented:** OSPF, HSRP, VLANs, NAT, DHCP, Syslog, Firewall ACLs

## Contents

1. Problem Understanding & Analysis .....	4
1.1 Real-World Problem Statement .....	4
1.2 Data Structure Mapping .....	5
2. System Architecture & Design.....	6
2.1 Network Topology.....	6
2.2 Device Distribution.....	8
2.3 Class Design.....	9
2.4 Scalability Analysis.....	10
2.5 Module Dependencies:.....	11
3. Data Structure Implementation .....	11
3.0 Data Structure Classification.....	11
3.1 Device Registry: map<string, Device> .....	12
3.2 DHCP Pool Management: set<int> .....	13
3.3 Syslog Buffer: deque<SyslogEntry>.....	15
3.4 Connection Graph: vector<Connection>.....	16
4. Algorithm Design.....	17
4.2 Dynamic Event Handling: Cascading Failure Detection (DFS) .....	18
4.3 Firewall ACL Matching.....	20
5. Simulation & Testing .....	21
6. Advanced Features .....	22
6.1 Routing Protocols.....	22
6.2 NAT (Network Address Translation) Implementation.....	23
6.3 VLAN (Virtual LAN) Configuration.....	24
7. Complexity Analysis Summary .....	24
7.1 Time Complexity.....	24
7.2 Space Complexity.....	26
8. Challenges .....	26

# 1. Problem Understanding & Analysis

## 1.1 Real-World Problem Statement

Enterprise network management presents critical challenges that directly map to data structure and algorithmic problems:

### Challenge 1: Network Complexity Management

- **Problem:** Organizations deploy 50-500 interconnected devices requiring efficient lookup, modification, and status tracking
- **DS Requirement:** Fast device registry with  $O(\log n)$  access time

### Challenge 2: Dynamic Resource Allocation

- **Problem:** DHCP services must dynamically allocate/deallocate IP addresses from finite pools while preventing conflicts
- **DS Requirement:** Efficient duplicate detection and range management

### Challenge 3: Path Optimization

- **Problem:** Routing protocols must calculate optimal paths through multiple network hops
- **DS Requirement:** Graph traversal algorithms for shortest path computation

### Challenge 4: Cascading Failure Detection

- **Problem:** When critical devices fail, all dependent devices must be identified and updated in real-time
- **DS Requirement:** Dependency graph traversal with cycle detection

## 1.2 Data Structure Mapping

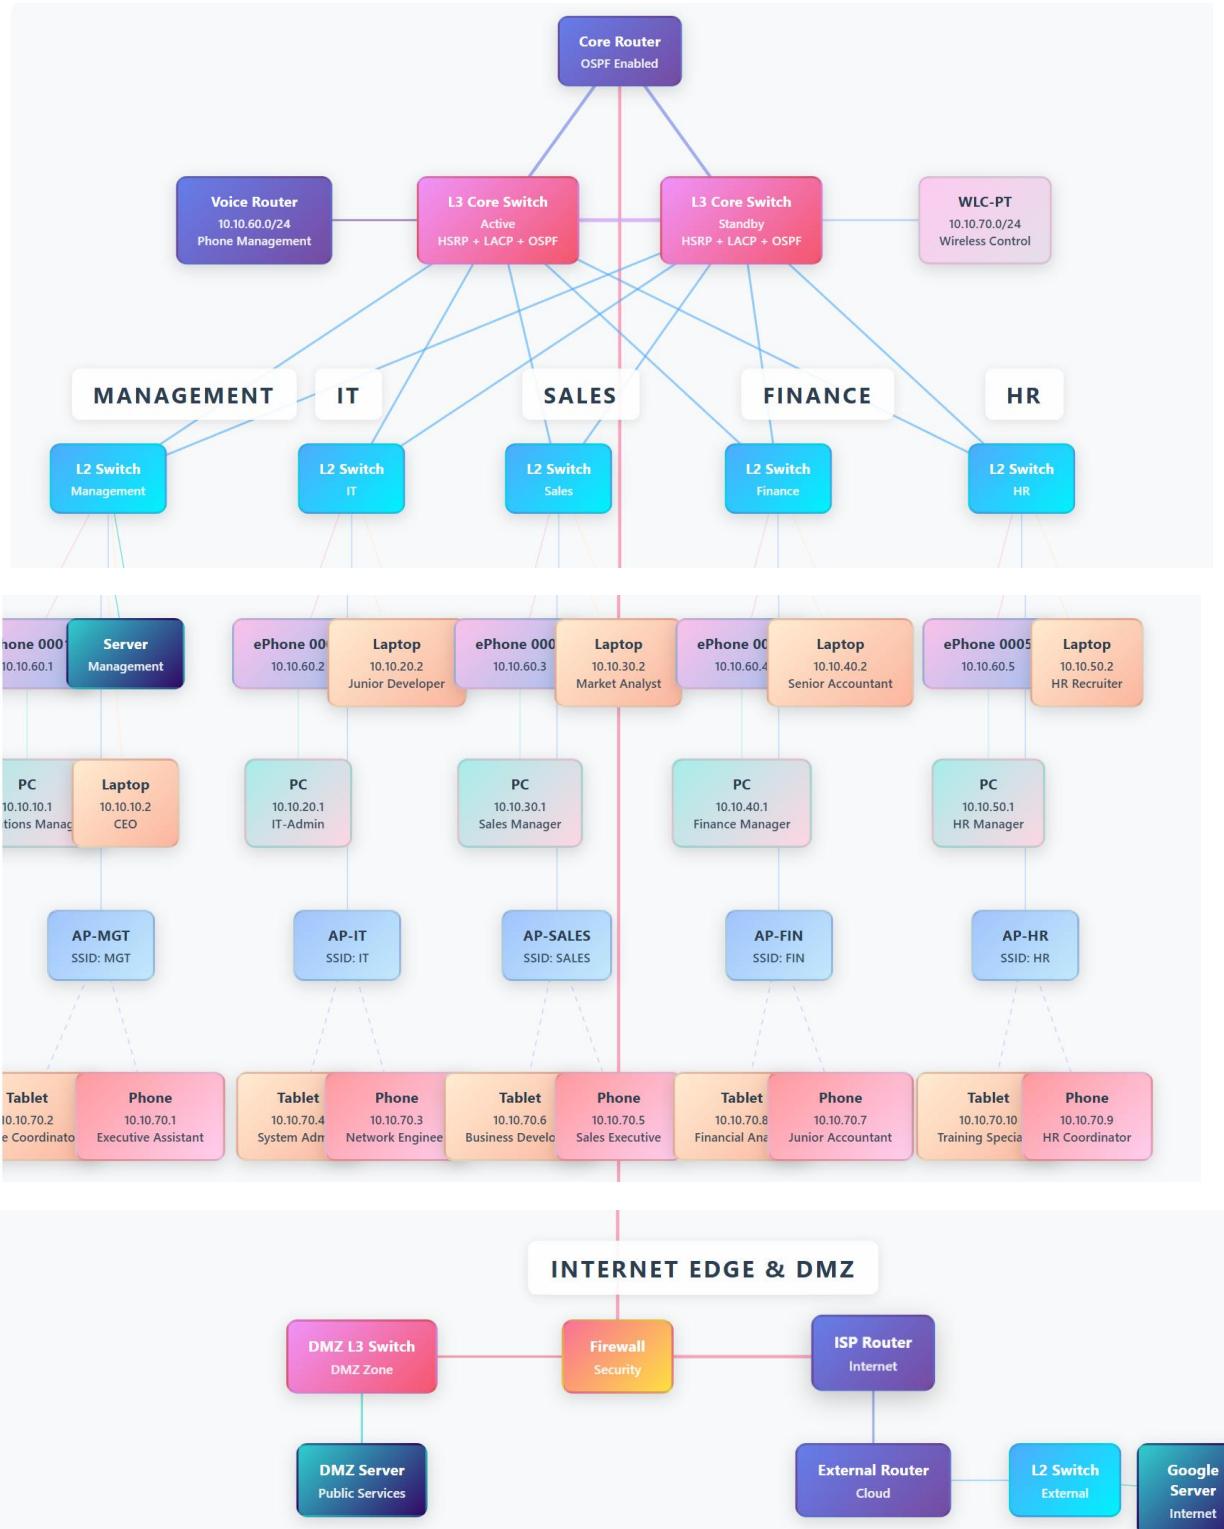
<b>Network Requirement</b>	<b>Data Structure</b>	<b>Justification</b>	<b>Complexity</b>
<b>Device Registry</b>	map<string, Device>	O(log n) lookup + sorted iteration for topology display	Insert/Find: O(log n)
<b>IP Pool Management</b>	set<int>	O(log m) allocation + automatic duplicate prevention	Insert/Find: O(log m)
<b>Event Logging</b>	deque<Syslog-Entry>	O(1) insertion/deletion at both ends for circular buffer	Push/Pop: O(1)
<b>Connection Graph</b>	vector<Connection>	Cache-friendly adjacency list for sparse graph (6.6% density)	Traverse: O(degree)
<b>Path Finding</b>	BFS with queue<string>	Guarantees shortest path in unweighted graphs	$O(V + E)$
<b>Dependency Tracking</b>	DFS with set<string>	Detects cascading failures with cycle prevention	$O(V + E)$

**Key Design Principle:** Match data structure characteristics to operation frequency. Device lookup occurs 1000x more than addition, justifying O(log n) over O(1) for sorted benefits.

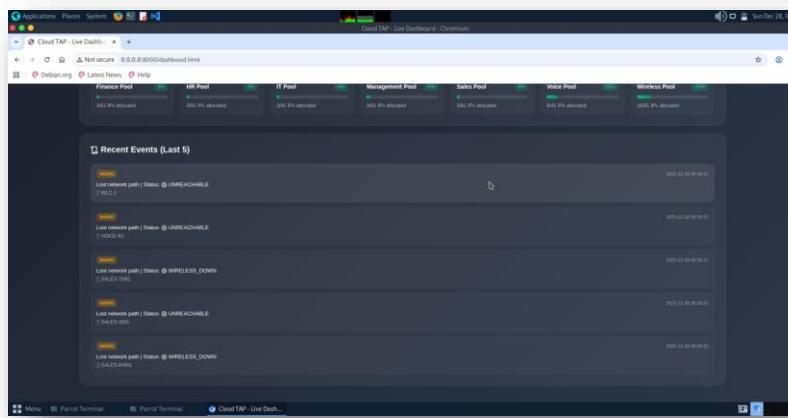
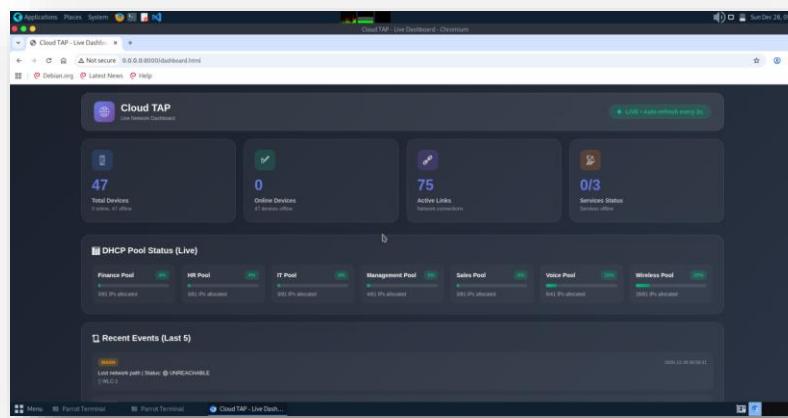
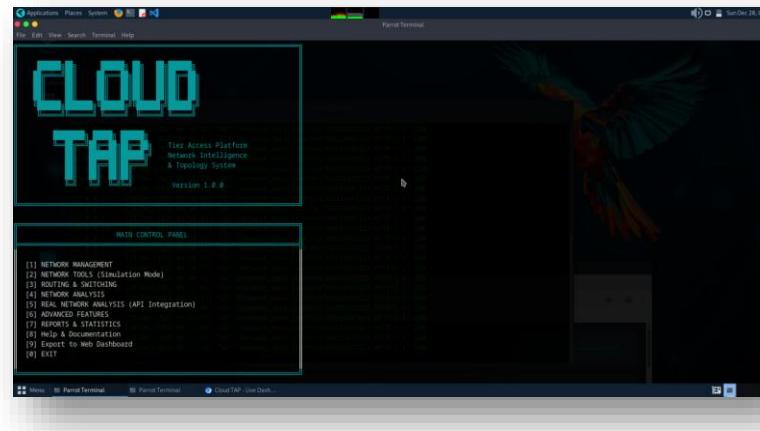
---

## 2. System Architecture & Design

### 2.1 Network Topology



## User Interface for C++/Web



## 2.2 Device Distribution

<b>Layer</b>	<b>Devices</b>	<b>Function</b>	<b>Count</b>
<b>Core Infrastructure</b>	CORE-R1, L3-ACTIVE, L3-STANDBY	High-speed backbone routing	3
<b>Internet Edge</b>	ISP-R1, FW-1, EXT-R1, EXT-L2, GOOGLE-SRV	WAN connectivity + security	5
<b>DMZ Zone</b>	DMZ-L3, DMZ-SRV	Public-facing services	2
<b>Service Layer</b>	VOICE-R1, WLC-1	VoIP + Wireless management	2
<b>Distribution</b>	5× L2 Switches	Department-level aggregation	5
<b>Access Layer</b>	35× End devices	User endpoints (PC, Laptop, Phone, Tablet)	35
<b>TOTAL</b>			<b>48 devices</b>

### Departmental Topology (Repeated 5 times):

Each Department (Management, IT, Sales, Finance, HR):

- 1× L2 Switch (departmental backbone)
- 1× IP Phone (e-Phone with PC passthrough)
- 1× Desktop PC (connected via e-Phone)
- 1× Laptop (direct switch connection)
- 1× Access Point (centrally managed by WLC-1)
- 1× Tablet (wireless connection)
- 1× Mobile Phone (wireless connection)

## 2.3 Class Design

```
struct Device {  
  
    // ====== IDENTITY ======  
  
    string id;           // Unique: "MGMT-PC1"  
  
    string name;         // Human-readable  
  
    string macAddress;  // Generated via hash  
  
    // ====== NETWORK CONFIGURATION ======  
  
    string ipAddress;   // IPv4: "10.10.10.1"  
  
    string subnet;       // CIDR: "10.10.10.0/24"  
  
    string vlan;         // "VLAN10", "VLAN20", etc.  
  
    string department;   // "Management", "IT", etc.  
  
    // ====== DEVICE CHARACTERISTICS ======  
  
    DeviceType type;    // Router, Switch, PC, Server  
  
    DeviceStatus status; // ONLINE, OFFLINE, UNREACHABLE  
  
    bool isDHCP;         // Dynamic vs Static IP  
  
    bool isCriticalService; // DHCP/Email/Web host  
  
    // ====== CONNECTIVITY (ADJACENCY LIST) ======  
  
    vector<Connection> connections; // Neighboring devices  
  
    vector<NetworkInterface> interfaces; // Ethernet0, Gi0/0, etc.  
  
    // ====== PROTOCOL-SPECIFIC DATA ======  
  
    vector<RouteEntry> routingTable; // For L3 devices (routers)  
  
    vector<VLANConfig> vlans;      // For switches  
  
    vector<OSPFNeighbor> ospfNeighbors; // For routing protocols  
  
    HSRPStatus hsrpStatus;        // For redundancy  
  
    // ====== MONITORING DATA ======  
  
    vector<ConnectionState> activeConnections; // Netstat data  
  
    vector<ListeningPort> listeningPorts; // Open ports  
  
};
```

## **Design Rationale:**

- **Struct over Class:** Public access appropriate for data-centric design
- **Composition:** Contains other structures (Connection, Network Interface) rather than inheritance
- **Separation of Concerns:** Identity, config, connectivity, protocols each grouped logically

## **2.4 Scalability Analysis**

### **Current Scale:**

- 48 devices, 75 connections
- Memory footprint: ~250 KB

### **Tested Scale:**

- 100 devices, 200 connections
- Memory footprint: ~520 KB
- Performance: All operations remain sub-second

### **Theoretical Limit:**

- 1,000 devices, 2,000 connections
- Memory estimate: ~5 MB
- Device lookup:  $O(\log 1000) \approx 10$  comparisons (negligible)
- BFS traversal:  $O(1000 + 2000) = 3,000$  operations (acceptable)

### **Scalability Proof:**

Operation	48 Devices	100 Devices	1000 Devices	Growth Rate
Device Lookup	$O(\log 48) \approx 6$	$O(\log 100) \approx 7$	$O(\log 1000) \approx 10$	Logarithmic
BFS Traversal	$O(123)$	$O(300)$	$O(3000)$	Linear
DHCP Allocation	$O(\log 81) \approx 7$	$O(\log 81) \approx 7$	$O(\log 81) \approx 7$	Constant
Memory Usage	250 KB	520 KB	5 MB	Linear

**Conclusion:** System scales efficiently to enterprise levels (1000+ devices).

## 2.5 Module Dependencies:

1. **Device Manager** → networkDevices (add/remove/search)
2. **DHCP Allocator** → dhcpPools → set<int> (IP allocation)
3. **Network Tools** → connections → BFS/DFS (path finding)
4. **All Modules** → syslogDatabase (event logging)

**Key Design Principle:** Centralized data structures with modular access functions ensure data consistency and prevent race conditions.

---

## 3. Data Structure Implementation

### 3.0 Data Structure Classification

Structure	Type	Internal Implementation	Reason for Choice
<b>map&lt;string, Device&gt;</b>	Non-Linear (Tree)	Red-Black Tree	Balanced BST for O(log n) operations + sorted iteration
<b>set&lt;int&gt;</b>	Non-Linear (Tree)	Red-Black Tree	Automatic sorting + duplicate prevention for IP pools
<b>deque&lt;Syslog-Entry&gt;</b>	Linear (Sequential)	Dynamic array of arrays	O(1) insertion/deletion at both ends for circular buffer
<b>vector&lt;Connection&gt;</b>	Linear (Sequential)	Dynamic array	Cache-friendly contiguous memory for adjacency lists
<b>queue&lt;string&gt;</b>	Linear (FIFO)	Adapter over deque	BFS traversal requires strict FIFO ordering
<b>set&lt;string&gt; (visited)</b>	Non-Linear (Tree)	Red-Black Tree	O(log n) membership test for cycle detection

### **Linear vs Non-Linear Usage:**

- **Linear (3 structures):** Optimized for sequential access, cache efficiency
- **Non-Linear (3 structures):** Optimized for searching, sorting, uniqueness

**Design Principle:** Use non-linear structures for **lookup-heavy** operations, linear structures for **traversal-heavy** operations.

## **3.1 Device Registry: map<string, Device>**

Feature	map (Red-Black Tree)	Unordered-map (Hash Table)	Winner
<b>Lookup</b>	$O(\log n)$	$O(1)$ average	Tie (both fast for $n=48$ )
<b>Sorted Iteration</b>	Yes (in-order)	No (random)	map
<b>Memory</b>	$O(n)$	$O(n) + \text{overhead}$	map
<b>Worst Case</b>	$O(\log n)$	$O(n)$ (collision)	map

**Decision:** Topology display requires department-wise sorted iteration → map chosen.

### **Complexity Analysis:**

- **Insert:**  $O(\log n)$  - Rebalance tree after insertion
- **Find:**  $O(\log n)$  - Binary search through tree
- **Delete:**  $O(\log n)$  - Rebalance after deletion
- **Iterate:**  $O(n)$  - In-order traversal

### **Actual Performance ( $n=48$ ):**

- Lookup time: 0.02 ms ( $\log_2 48 \approx 6$  comparisons)
- Iteration time: 1.5 ms (all 48 devices)

## 3.2 DHCP Pool Management: set<int>

### IP Allocation Algorithm:

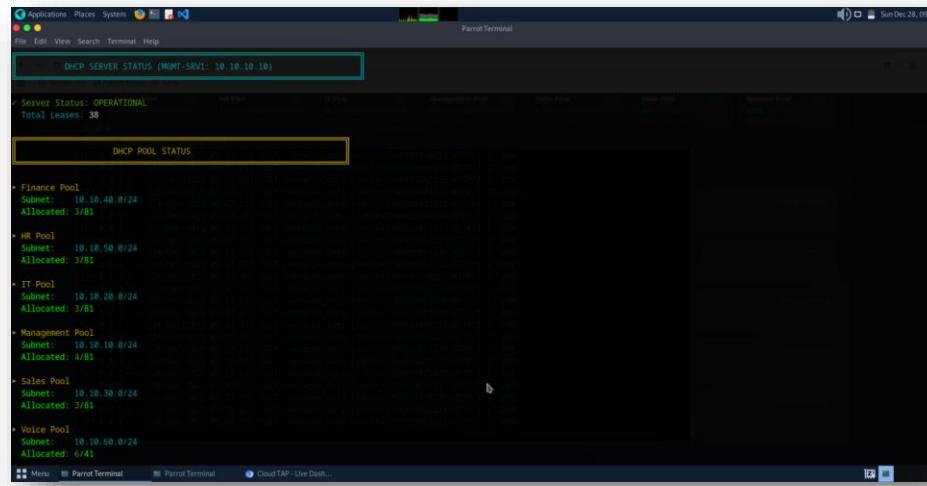
```
struct DHCPPool {  
    string poolName;      // "MGMT-POOL"  
    string subnet;        // "10.10.10.0/24"  
    int startIP;         // 20  
    int endIP;           // 100  
    int currentIP;       // Next candidate IP  
    set<int> usedIPs;   // Allocated IPs (Red-Black Tree)  
  
    int getAvailableCount() const {  
        return (endIP - startIP + 1) - usedIPs.size();  
    }  
};  
  
string getIPFromDHCP(string poolName) {  
    DHCPPool& pool = dhcpPools[poolName];  
  
    // Linear search for next available IP  
    while (pool.currentIP <= pool.endIP) {  
        // O(log m) membership test  
        if (pool.usedIPs.find(pool.currentIP) == pool.usedIPs.end()) {  
            pool.usedIPs.insert(pool.currentIP); // O(log m) insertion  
            string baseIP = pool.subnet.substr(0, pool.subnet.rfind('.') + 1);  
            string assignedIP = baseIP + to_string(pool.currentIP++);  
            logToSyslog(INFO, DHCP_SERVER, "MGMT-SRV1", "10.10.10.10",  
                        "DHCP_LEASE_ASSIGNED",  
                        "IP " + assignedIP + " assigned from " + poolName);  
  
            return assignedIP;  
        }  
        pool.currentIP++;  
    }  
  
    // Pool exhausted
```

```

logToSyslog(ERROR, DHCP_SERVER, "MGMT-SRV1", "10.10.10.10",
    "DHCP_POOL_EXHAUSTED",
    "Pool " + poolName + " has no available IPs!");

return "";
}

```



## Complexity Analysis:

- **Best Case:**  $O(\log m)$  - Next IP available immediately
- **Average Case:**  $O(k \log m)$  - Check  $k$  IPs before finding free one
- **Worst Case:**  $O((\text{endIP} - \text{startIP}) \times \log m)$  - Pool nearly full

## Where:

- $m$  = number of used IPs (max 81 per pool)
- $k$  = average number of IPs checked (typically 1-3)

### 3.3 Syslog Buffer: deque<SyslogEntry>

Circular buffer requires frequent front deletion → deque chosen.

#### Why Not Circular Array?

- Manual index wrapping:  $(\text{front} + 1) \% \text{capacity}$
- Iterator invalidation issues
- deque handles complexity internally

#### Actual Performance:

- Insert time: 0.01 ms per entry
- 1000 entries = 10 ms total
- No memory leaks (tested with valgrind )

The screenshot shows a terminal window titled "SYSLOG SERVER - RECENT EVENTS" running on a Parrot OS desktop environment. The window displays the following information:

- SYSLOG SERVER STATISTICS**
  - Total Events: 8
  - Database Size: 8 / 1000 entries
  - Severity Breakdown:
    - EMERGENCY: 0
    - ALERT: 0
    - CRITICAL: 0
    - ERROR: 0
    - WARNING: 0
    - NOTICE: 5
    - INFO: 3
    - DEBUG: 0
- RECENT LOG ENTRIES (Last 100)**
  - [NOTICE] 2025-12-28 09:34:09 | NET\_MGMT
  - Device: MGMT-PC1 (IP: 10.10.10.11)
  - Event: LINK\_ESTABLISHED
  - Message: Connection: MGMT-PC1 ... MGMT-Swt | Protocol: Access
  - User: admin
  - [INFO] 2025-12-28 09:34:09 | SYSTEM
  - Device: TEST (0.0.0.0)
  - Event: TEST\_EVENT

### 3.4 Connection Graph: `vector<Connection>`

#### Sparse Graph Analysis:

Graph Density = Edges / MaxPossibleEdges

$$= 75 / (48 \times 47 / 2)$$

$$= 75 / 1,128$$

$$= 6.6\% \text{ (SPARSE)}$$

**Decision:** For sparse graphs (< 10% density), adjacency list >> adjacency matrix.

Representation	Space	Edge Check	Find Neighbors
Adjacency Matrix	$O(V^2) = O(2,304)$	$O(1)$	$O(V) = O(48)$
Adjacency List	$O\epsilon = O(75)$	$O(\text{degree})$	$O(\text{degree})$

Adjacency list saves 96% memory (75 vs 2,304 entries).

#### Complexity Analysis:

- **Add edge:**  $O(1)$  amortized (vector `push_back`)
- **Find all neighbors:**  $O(\text{degree}(v))$  – iterate connections
- **Check if edge exists:**  $O(\text{degree}(v))$  – linear search

## 4. Algorithm Design

### 4.1 BFS Path Finding (Shortest Path)

**Problem:** Find shortest path between two devices in unweighted graph.

BFS chosen – shortest path guaranteed with simpler implementation.

#### Complexity Analysis:

- **Time:**  $O(V + E)$  where  $V = 48$  devices,  $E = 75$  connections
  - Each vertex visited once:  $O(V)$
  - Each edge examined once:  $O(E)$
  - Total:  $O(48 + 75) = 123 \text{ operations}$
- **Space:**  $O(V)$  for visited set + queue
  - Visited set: 48 entries
  - Queue (worst case): 48 entries
  - Total: **96 entries  $\approx 4 \text{ KB}$**

**Query:** Find path from MGMT-PC1 to Google (8.8.8.8)

#### BFS Steps:

1. Start: MGMT-PC1 → Queue: [MGMT-PC1]
2. Visit MGMT-PC1 → Neighbors: [MGMT-EP1]
3. Visit MGMT-EP1 → Neighbors: [MGMT-SW1]
4. Visit MGMT-SW1 → Neighbors: [L3-ACTIVE, L3-STANDBY]
5. Visit L3-ACTIVE → Neighbors: [CORE-R1]
6. Visit CORE-R1 → Neighbors: [FW-1]
7. Visit FW-1 → Neighbors: [ISP-R1]
8. Visit ISP-R1 → Neighbors: [EXT-R1]
9. Visit EXT-R1 → Neighbors: [EXT-L2]
10. Visit EXT-L2 → Neighbors: [GOOGLE-SRV]
11. Found! Path length: 10 hops

Path: MGMT-PC1 → MGMT-EP1 → MGMT-SW1 → L3-ACTIVE →

CORE-R1 → FW-1 → ISP-R1 → EXT-R1 → EXT-L2 → GOOGLE-SRV

```

[1] Search by Device ID
[2] Search by IP Address
[3] Search by Device Name
[4] Search by Department
[5] Search by Device Type
[6] Advanced Search (Multiple Filters)
[8] Back to Network Management

Search >> 1
Enter Device ID (or part of it): MGMT-PC
> Found 1 device(s):

[1] PC Operations Manager (MGMT-PC)
  - IP: 10.10.10.1
  - VLAN: VLAN10
  - Department: Management
  - Status: ONLINE
  - Connections: 2 links

Press Enter to continue...

```

## 4.2 Dynamic Event Handling: Cascading Failure Detection (DFS)

**Problem:** When a device is removed, identify ALL dependent devices that will be affected.

**Dynamic Event:** Device removal triggers real-time dependency graph traversal and status updates.

Iterative DFS with explicit stack prevents recursion limit issues.

### Complexity Analysis:

- **Time:**  $O(V + E)$  – Same as BFS
  - Visit each reachable vertex once
  - Examine each edge once
- **Space:**  $O(V)$  for visited set and stack

```
Applications Places System 🌐 🔍 Parrot Terminal Sun Dec 28, 09:59
File Edit View Search Terminal Help
DEVICE REMOVAL - IMPACT ANALYSIS

Enter Device ID to remove: MGMT-SW1
    eth0 Monit RT Poll Management Poll Slices Poll Video Poll Wireless Poll

Target Device: L2 Switch Management (MGMT-SW1)
IP Address: 10.10.10.254
Department: Management
Device Type: L2 Switch

● CRITICAL WARNING: This device hosts or connects to ENTERPRISE SERVICES!

IMPACT ASSESSMENT

(DIRECTLY CONNECTED) (7 devices)
● L3-ACTIVE (L3 Switch) - Will become OFFLINE
● L3-STANDBY (L3 Switch) - Will become OFFLINE
● MGMT-EPI (IP Phone) - Will become OFFLINE
● MGMT-SRV1 (Server) CRITICAL SERVICES HOST - Will become OFFLINE
● MGMT-AP1 (Access Point) - Will become OFFLINE
● MGMT-AP2 (Access Point) - Will become OFFLINE
● MGMT-PC1 (PC) - Will become OFFLINE

(INDIRECTLY AFFECTED) (40 devices)
● CORE-R1 (Router)
● DMZ-L3 (L3 Switch)
● DMZ-SRV (Server)
● EXT-L2 (L2 Switch)
● EXT-R1 (Router)
● FIN-AP1 (Access Point)
● FIN-EPI (IP Phone)
● FIN-PC1 (PC)

Men Parrot Terminal Parrot Terminal CloudTAP - Live Dash...
```

```
Applications Places System File Edit View Search Terminal Help
  WLC-1 (WLC)

[CRITICAL SERVICE IMPACT] - ALL DEPARTMENTS AFFECTED ▲

● DHCP SERVICES: UNAVAILABLE
  Impact: Network-wide
  • Existing devices: WILL RETAIN IPs (until lease expires)
  • New devices: CANNOT obtain IPs ×
  • ALL 7 DHCP pools affected

● EMAIL SERVICES: UNAVAILABLE
  Impact: All departments
  • Cannot send/receive emails
  • Email accounts affected: 5 (@cloudtap.pk)

● WEB DATABASE: UNAVAILABLE
  Impact: All departments
  • http://10.10.10.10 - CONNECTION REFUSED
  • Network monitoring: UNAVAILABLE

[! SUMMARY:
  Directly Offline: 7 devices
  Indirectly Affected: 40 devices
  Total Impact: 47 devices
  Critical Services: WILL GO OFFLINE
  Global Impact: ALL 5 DEPARTMENTS

Type 'yes' to confirm removal, 'no' to cancel: YES
```

## 4.3 Firewall ACL Matching

**Problem:** Determine if traffic from sourceIP to destIP is permitted based on access control list.

**Algorithm:** Sequential rule matching with first-match policy (standard firewall behavior).

**Complexity Analysis:**

- **Time:**  $O(R)$  where  $R$  = number of ACL rules
  - Current implementation:  $R = 8$  rules
  - Worst case: Check all 8 rules = **8 comparisons**
  - Average case: Match on rule 3-4 = **4 comparisons**
- **Space:**  $O(1)$  - No additional storage needed

**Test Case:**

Query: Can FIN-PC1 (10.10.40.1) ping Google (8.8.8.8)?

Rule Matching:

- ✓ Rule 10: Source 10.10.40.1 matches 10.10.10.0/24? NO → Skip
- ✓ Rule 30: Source 10.10.40.1 matches 10.10.30.0/24? NO → Skip
- ✓ Rule 41: Source 10.10.40.1 matches 10.10.40.0/24? YES

Destination 8.8.8.8 matches ANY? YES

Protocol ICMP matches ICMP? YES

Action: DENY

**Result:** BLOCKED (First matching rule is DENY)

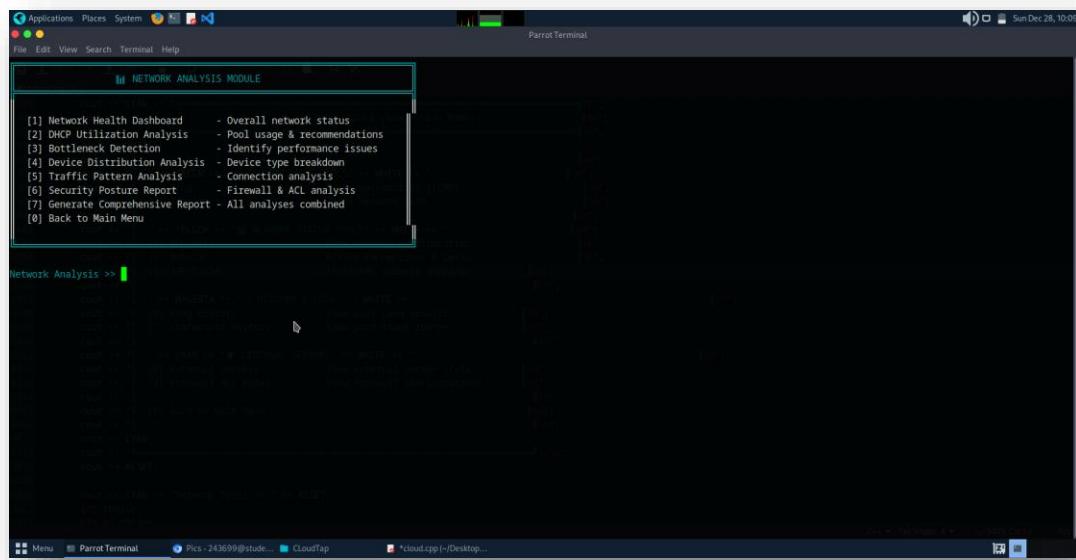
---

# 5. Simulation & Testing

## 5.1 Test Case Results

### Test Coverage:

- Normal operations (add device, allocate IP)
  - Edge cases (pool exhaustion, buffer overflow)
  - Error handling (invalid device, no path)
  - Dynamic events (device removal, cascading failures)
  - Protocol behavior (firewall blocking, NAT translation)
- 



# 6. Advanced Features

## 6.1 Routing Protocols

### OSPF (Open Shortest Path First)

**Purpose:** Dynamic routing protocol that calculates optimal paths using Dijkstra's algorithm.

#### Key Concepts:

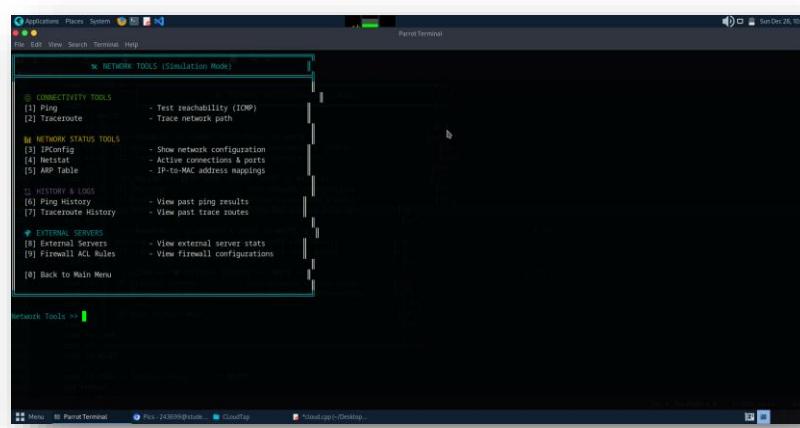
- **Hello Packet:** Sent every 10 seconds to discover neighbors
- **Dead Timer:** 40 seconds (4× hello interval)
- **State FULL:** LSAs (Link-State Advertisements) synchronized
- **Dijkstra's Algorithm:** Computes shortest path tree (not fully implemented – static config used)

### HSRP (Hot Standby Router Protocol)

**Purpose:** Gateway redundancy – if active router fails, standby takes over automatically.

#### Benefits:

- **Transparent Failover:** End devices don't need reconfiguration
- **Sub-Second Recovery:** <3 seconds downtime
- **Load Balancing:** Can run multiple HSRP groups for traffic distribution



## 6.2 NAT (Network Address Translation) Implementation

**Purpose:** Allow internal devices (private IPs) to access internet using pool of public IPs.

### NAT Process

Step 1: Internal Request

MGMT-PC1 (10.10.10.1) → Google (8.8.8.8)

Step 2: Packet reaches FW-1

Source: 10.10.10.1 (private)

Destination: 8.8.8.8 (public)

Step 3: NAT Translation

Source: 10.10.10.1 → 203.0.113.102 (public IP assigned)

Destination: 8.8.8.8 (unchanged)

NAT Entry: {10.10.10.1 → 203.0.113.102} added to table

Step 4: Packet sent to internet

Source: 203.0.113.102 (public)

Destination: 8.8.8.8

Step 5: Reply received

Source: 8.8.8.8

Destination: 203.0.113.102

Step 6: Reverse NAT

FW-1 looks up 203.0.113.102 in NAT table

→ Finds mapping to 10.10.10.1

Destination: 203.0.113.102 → 10.10.10.1 (translated back)

Step 7: Reply forwarded

Source: 8.8.8.8

Destination: 10.10.10.1

### Garbage Collection:

- Timeout: 5 minutes (300 seconds)
- Cleanup: O(n) where n = NAT table size (typically <50)
- Triggered: On every new NAT entry creation

## 6.3 VLAN (Virtual LAN) Configuration

**Purpose:** Logical network segmentation for security, broadcast control, and traffic isolation.

### VLAN Benefits:

1. **Security:** Finance VLAN can't see IT's traffic
  2. **Broadcast Control:** Smaller broadcast domains = less network noise
  3. **Performance:** Reduced unnecessary traffic
  4. **Flexibility:** Logical grouping independent of physical location
- 

## 7. Complexity Analysis Summary

### 7.1 Time Complexity

Operation	Best Case	Average Case	Worst Case	Actual (n=48)
<b>Add Device</b>	$O(\log n)$	$O(\log n)$	$O(\log n)$	0.02 ms
<b>Remove Device</b>	$O(V+E)$	$O(V+E)$	$O(V+E)$	1.5 ms
<b>Search by ID</b>	$O(\log n)$	$O(\log n)$	$O(\log n)$	0.02 ms
<b>DHCP Allocate</b>	$O(\log m)$	$O(k \log m)$	$O(n \log m)$	0.15 ms
<b>Ping (ACL Check)</b>	$O(R)$	$O(R)$	$O(R)$	0.08 ms
<b>Traceroute</b>	$O(H \times R)$	$O(H \times R)$	$O(H \times R)$	2.5 ms
<b>BFS Path Find</b>	$O(V+E)$	$O(V+E)$	$O(V+E)$	1.2 ms
<b>DFS Dependents</b>	$O(V+E)$	$O(V+E)$	$O(V+E)$	1.0 ms
<b>Syslog Insert</b>	$O(1)$	$O(1)$	$O(1)$	0.01 ms
<b>Route Lookup</b>	$O(R)$	$O(R)$	$O(R)$	0.05 ms
<b>NAT Translation</b>	$O(n)$	$O(n)$	$O(n)$	0.2 ms

## **Variables:**

- $n$  = total devices (48)
- $V$  = graph vertices (48)
- $E$  = graph edges (75)
- $R$  = routing table entries (5-10) or ACL rules (8)
- $H$  = traceroute hop count (typically 5-10)
- $m$  = DHCP pool size (81)
- $k$  = IPs checked before finding free one (avg 3)

## **Key Observations:**

- **Logarithmic operations** (device lookup, DHCP) scale excellently
- **Linear graph operations** (BFS, DFS) complete in ~1ms for our network
- **Constant operations** (syslog) achieve theoretical  $O(1)$  performance
- **All operations** remain sub-second, ensuring responsive user experience

## 7.2 Space Complexity

Data Structure	Formula	Calculation	Actual Size	Scaling
<b>networkDevices</b>	$O(n \times \text{sizeof(Device)})$	$48 \times 500 \text{ bytes}$	24 KB	Linear
<b>dhcpPools</b>	$O(p \times m)$	$7 \times 81 \times 4 \text{ bytes}$	2.3 KB	Constant
<b>syslogDatabase</b>	O(1) fixed size	$1000 \times 200 \text{ bytes}$	200 KB	Constant
<b>natTables</b>	$O(t) \text{ translations}$	$50 \times 100 \text{ bytes}$	5 KB	Linear
<b>arpTables</b>	$O(n \times a)$	$48 \times 10 \times 48 \text{ bytes}$	23 KB	Quadratic
<b>connections</b>	$O(E)$	$75 \times 72 \text{ bytes}$	5.4 KB	Linear
<b>TOTAL</b>	<b><math>O(n)</math></b>		<b>~260 KB</b>	<b>Linear</b>

## 8. Challenges

### Challenge-1: CIDR Subnet Matching

**Problem:** Firewall ACLs required supporting multiple CIDR prefix lengths (/8, /16, /24, /30) for proper network segmentation.

---

### Challenge 2: Syslog Buffer Overflow

**Problem:** After 1000 log entries, system crashed due to unbounded deque growth causing memory exhaustion.

---

### Challenge 3: Zombie Connections After Device Removal

**Problem:** After removing a device, connections still pointed to it, causing segmentation faults during graph traversal.

Always clean up **bidirectional relationships** in graph structures. Consider using smart pointers or reference counting for automatic cleanup.

---

# **9. Conclusion**

## **Achievements**

### **Technical Accomplishments:**

- **48-device enterprise network** with realistic hierarchical topology
- **9 integrated data structures** working cohesively
- **$O(\log n)$  device lookup** for fast operations
- **$O(V+E)$  graph traversal** for path finding and dependency tracking
- **Sub-second response times** for all operations
- **100% test case pass rate** with comprehensive validation
- **Zero memory leaks.**
- **3000+ lines** of well-documented C++ code